Rohan Goel, Paru Swaminathan, Satej Shah

1. Progress made thus far

Thus far, we have set up our working environments and installed the necessary software. We have understood the task at hand and created a plan to move forward. We have cleaned the training data by removing stop words and brackets and by converting the data into word2vec vectors.

2. Remaining tasks

We will now be building a model to learn to classify whether the response with context is sarcastic or not. We are planning on trying out Naive Bayes, Linear Regression, RNNs, a simple Neural Network, and a simple Random Forest Classifier. We plan to test each of these models in the week of November 30th, and are hoping to move forward with one of them, and begin hyperparameter tuning by December 7th.

3. Any challenges/issues being faced

Currently, we are facing issues with how to utilize both the response and context. There are many tutorials online that help with understanding how to use one column, however, working with multiple features proves to be more challenging.